Processing of Probabilistic Skyline Queries Using MapReduce

Qingfu Wen

School of Software, Tsinghua University

qingfu.wen@gmail.com

Author: Yoonjae Park, Jun-Ki Min, Kyuseok Shim

November 16, 2015

Probabilistic Skylines

Object	Instance	NO_2	SO_2	Probability
W	w_1	10	40	0.5
	w_2	75	10	0.4
X	x_1	55	20	0.2
	x_2	65	30	0.2
Y	y_1	95	60	0.8
	y_2	80	70	0.2
Z	z_1	5	80	0.5
	z_2	90	25	0.5

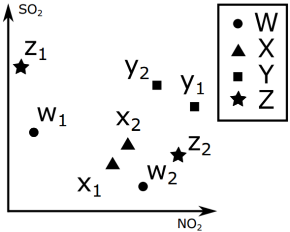

Probabilistic Skylines

Object	Instance	NO_2	SO_2	Probability
W	w_1	10	40	0.5
	w_2	75	10	0.4
X	x_1	55	20	0.2
	x_2	65	30	0.2
Y	y_1	95	60	0.8
	y_2	80	70	0.2
Z	z_1	5	80	0.5
	z_2	90	25	0.5

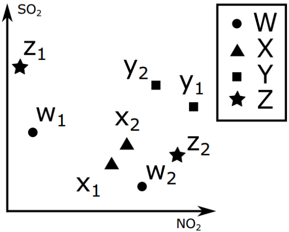

question

Which Object's skyline probability is larger than 0.6?

Probabilistic Skylines

Object	Instance	NO_2	SO_2	Probability
W	w_1	10	40	0.5
	w_2	75	10	0.4
X	x_1	55	20	0.2
	x_2	65	30	0.2
Y	y_1	95	60	0.8
	y_2	80	70	0.2
Z	z_1	5	80	0.5
	z_2	90	25	0.5

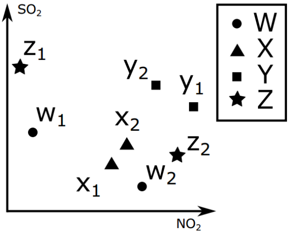

$$
\begin{aligned}
P_{sky}(y_1) &= P(y_1)(1 - P(w_1) - P(w_2))(1 - P(x_1) - P(x_2))(1 - P(z_2)) \\
&= 0.024 \\
P_{sky}(y_2) &= 0.012 \\
P_{sky}(Y) &= P_{sky}(y_1) + P_{sky}(y_2) = 0.036 \\
P_{sky}(W) &= 0.9 \\
P_{sky}(X) &= 0.4 \\
P_{sky}(Z) &= 0.74
\end{aligned}
$$

Probabilistic Skylines

Probabilistic Skyline Problem

For a set of uncertain objects $\mathbb{D}$ and a probability threshold T_p, the probabilistic skyline $pSL(\mathbb{D}, T_p)$, is the set of all objects whose skyline probabilities are at least T_p, $pSL(\mathbb{D}, T_p) = \{U \in \mathbb{D} | P_{sky}(U) \geq T_p\}$.

The discrete model:

$$P_{sky}(U) = \sum_{u_i \in U} P_{sky}(u_i) = \sum_{u_i \in U} (P(u_i) \times \prod_{V \in \mathbb{D}, V \neq U} (1 - \sum_{v_j \in V, v_j \prec u_i} P(v_j)))$$

The continuous model:

$$P_{sky}(u_i) = \int_U Uf(u) \times \prod_{V \in \mathbb{D}, V \neq U} (1 - \int_V Vf(v) 1(v \prec u)\, dv)\, du$$

What is MapReduce?

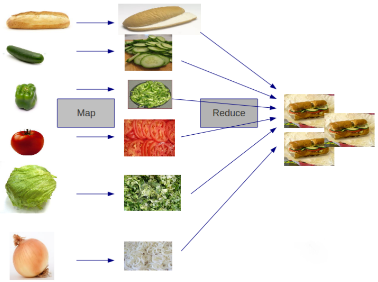

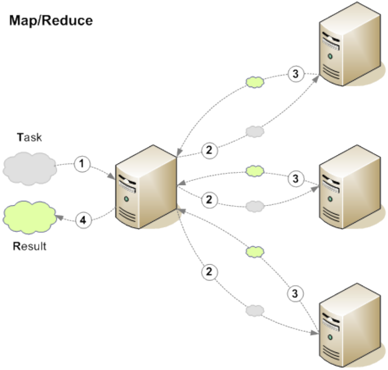

PSMR: The State-of-the-art Algorithm

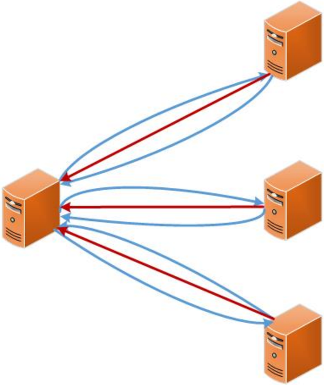

1. local computing candidate sets.
2. merge candidate sets, broadcast and local computing, reduce probabilities.

Early Pruning Techniques

Lemma (Zero-probability Filtering)

$$P_{sky}(U) = \sum_{u_i \in U} P(u_i) \times \prod_{V \in \mathbb{D}, V \neq U} (1 - \sum_{v_j \in V, v_j \prec u_i} P(v_j))$$

$$delete\ u_i\ if \prod_{V \in \mathbb{D}, V \neq U} (1 - \sum_{v_j \in V, v_j \prec u_i} P(v_j)) = 0.$$

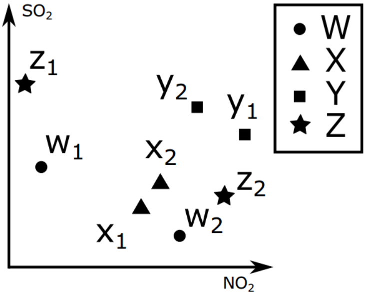

Example (Zero-probability Filtering)

we can see, x_1 dominate y_1, if $P(x_1) = 1$, then $P_{sky}(y_1) = 0$

Early Pruning Techniques

Lemma (Upper-bound Filtering)

$$\beta(U, \mathbb{S}, R(u_i)) = \frac{\prod_{V \in \mathbb{S}} (1 - \sum_{v_j \in V, v_j \prec R(u_i).min} P(v_j))}{1 - \sum_{v_k \in U, v_k \prec R(u_k).min} P(v_j)}$$

$$up(u_i, U, \mathbb{S}, R(u_i)) = P(u_i) \times \beta(U, \mathbb{S}, R(u_i)).$$

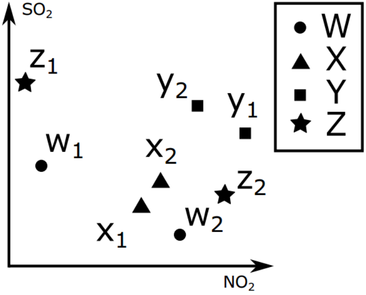

Example (Upper-bound Filtering)

$$P_{sky}(y_1) = P(y_1)(1 - P(w_1) - P(w_2))(1 - P(x_1) - P(x_2))(1 - P(z_2))$$
$$\leq P(y_1)(1 - P(w_1)) = 0.4$$
$$P_{sky}(y_2) \leq 0.1$$
$$P_{sky}(Y) = P_{sky}(y_1) + P_{sky}(y_2) \leq 0.5 < T_p$$

Early Pruning Techniques

Lemma (Dominance-Power Filtering)

$DP(v_j) = \prod_{1}^{d}(b(k) - v_j(k)) = 0, b(k) = max\{v_1(k), \cdots, v_n(k)\}.$

$DP(V) = \sum_{v_j \in V}(P(v_j) \times DP(v_j)).$

$\mathbb{F}$ *is topK DP set,* $\sum\limits_{u_i \in U} P(u_i) \times \prod\limits_{V \in \mathbb{F}, V \neq U} (1 - \sum\limits_{v_j \in V, v_j \prec u_i} P(v_j)) < T_p$*, U is not a probabilistic skyline Object.*

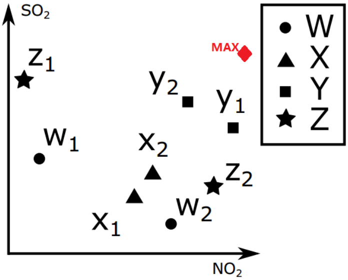

Example (Dominance-Power Filtering)

$$DP(y_1) = (100 - 95)(100 - 60) = 200$$
$$DP(y_2) = (100 - 80)(100 - 70) = 600$$
$$DP(Y) = P(y_1) * DP(y_1) + P(y_2) * DP(y_2)$$
$$= 0.8 * 200 + 0.2 * 600 = 280$$

PSQtree for Pruning

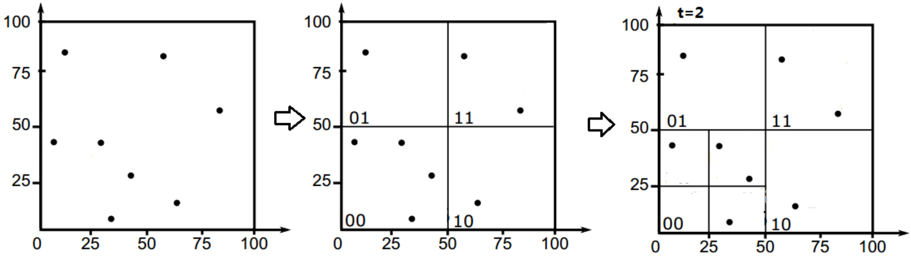

- generate PSQtree
- traverse PSQtree for computing $P_{sky}(node.min)$
- zero-probability filtering
- upper-bound filtering
- partitioning objects by PSQtree(weakly dominate)

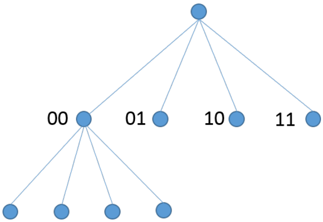

MapReduce Algorithms with PSQtree

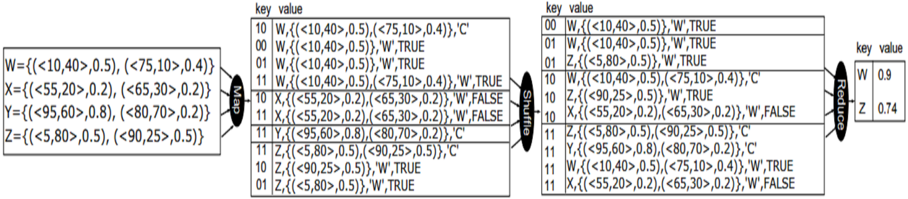

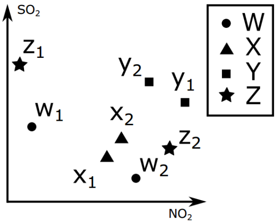

Example (PS-QPF-MR)

- $T_p = 0.5$, map function is called with an uncertain object.
- for X, upper bound $node(10).P_{min}P(x1) + node(10).P_{min}P(x2) = 0.4 < T_p = 0.5$. X is not a skyline candidate.
- every instance of X, emit the key-value pairs $\langle 10, (\{(\langle 55, 20\rangle, 0.2), (\langle 65, 30\rangle, 0.2)\}, W, False)\rangle$ *and* $\langle 11, (\{(\langle 55, 20\rangle, 0.2), (\langle 65, 30\rangle, 0.2)\}, W, False)\rangle$ since node(10) weakly dominates node(10) and node(11)

Experiments

- 50 machines with Intel i3 3.3GHz CPU and 4GB, Linux
- 200 machines with Intel Xeon 2.5GHz CPU and 3.75GB, Amazon EC2
- Java 1.6, Hadoop 1.2.1

Algorithm	Description
PS-QP-MR	The algorithm with quadtree partitioning
PS-QPF-MR	The algorithm with quadtree partitioning and filtering
PS-BR-MR	The algorithm with random partitioning
PS-BRF-MR	The algorithm with random partitioning and filtering
PSMR	The state-of-the-art algorithm

Parameter	Range	Default
No. of samples($\mathbb{S}$)	1000$\sim$10,000	1000 for PS-QPF-MR 2000 for PS-QP-MR 10000 for PS-BRF-MR
No. of dominating objects($\mathbb{F}$)	1000$\sim$10,000	100 for PS-QPF-MR 1000 for PS-BRF-MR
No. of objects($\mathbb{D}$)	$10^5 \sim 10^8$	10^7
No. of dimensions(d)	$2 \sim 8$	4
Probability threshold (T_p)	0.1$\sim$0.6	0.3
No. of inst. per object(ℓ)	$1 \sim 400$	40
No. of machines (t)	10$\sim$200	25

Experiments

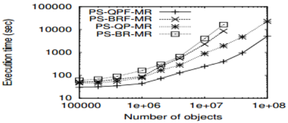

(a) ANTI

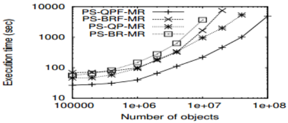

(b) IND

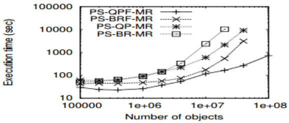

(c) COR

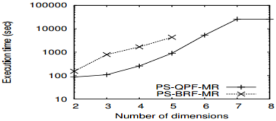

(a) ANTI

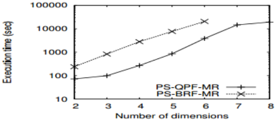

(b) IND

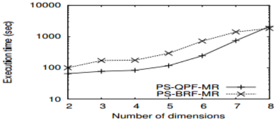

(c) COR

Experiments

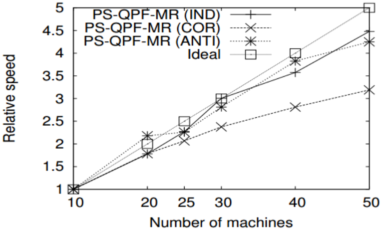

(a) With our cluster

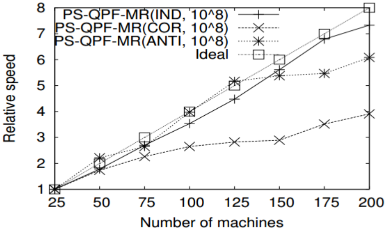

(b) With Amazon EC2

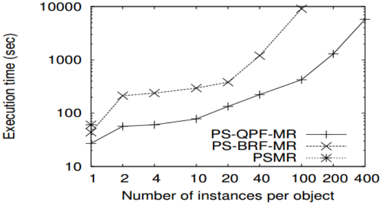

(a) Varying ℓ

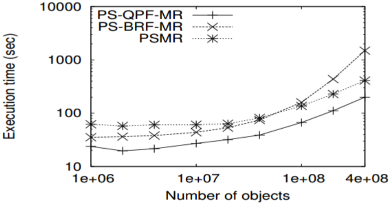

(b) Varying $|\mathbb{D}|$ when $\ell = 1$

Conclusion

- probabilistic skyline query for both discrete and continuous models
- zero-probability, the upper-bound, and dominance power filtering techniques
- using a PSQtree to distribute the instances of objects effectively
- a single MapReduce phase algorithm PS-QPF-MR and grouping techniques for optimization

The End

Q & A